

CSE 537 Assignment 1 Report: Peg Solitaire

Remy Oukaour
SBU ID: 107122849
remy.oukaour@gmail.com

Jian Yang
SBU ID: 110168771
jian.yang.1@stonybrook.edu

Thursday, September 24, 2015

1 Introduction

This report describes our submission for assignment 1 in the CSE 537 course on artificial intelligence. Assignment 1 requires us to implement three algorithms for solving a game of Peg Solitaire: iterative-deepening depth-first search (ID-DFS), and A* search with two different heuristics. In this report, we discuss the implementation details and compare the performance of six solutions (ID-DFS and five A* heuristics, of which we settled on two).

2 Implementation

We were given four Python code files implementing the assignment framework: *readGame.py*, *pegSolitaireUtils.py*, *search.py*, and *pegSolitaire.py*. We did not need to edit *readGame.py* or *pegSolitaire.py*, but only had to implement utility classes and methods in *pegSolitaireUtils.py* and the search algorithms in *search.py*.

2.1 Data structures

The original data structure given in *pegSolitaireUtils.py* was a single class *game* containing the game state as a 2-dimensional 7×7 array of integers, a trace of the moves taken to reach that state, and an overall count of the nodes expanded by the search algorithm. To make the implementation more clear, we split this into two classes: *game*, a singleton class storing overall state for a search (the original game state, final move trace, and expanded node count), and *gameNode*, a class which a search algorithm can repeatedly create as it explores the game tree.

The *gameNode* class contains some optimizations. We flattened the game state into a 1-dimensional array of 49 integers to speed up indexing and simplify copying. Thus, an index *gameState*[*i*][*j*] would become *gameState*[*i* * 7 + *j*]. (We did not change the state representation of *game*, only of *gameNode*.) We also cache two pieces of data when a *gameNode* is created: a count of pegs remaining on the board, and a *key* state chosen from the eight symmetric states under rotation and reflection. This data is used by our search algorithms to prune search trees.

2.2 Utility methods

We were required to implement four utility methods in *pegSolitaireUtils.py*:

1. *is_corner(self, pos)*:

```
def is_corner(self, pos):  
    """  
    Return whether a position is a corner/wall,  
    i.e. not a hole or peg.
```

```

"""
return 0 <= pos < 49 and self.state[pos] != -1

```

In implementation, we actually used another function `__getitem__(self, pos)` instead:

```

def __getitem__(self, pos):
    """
    node[pos] is bounds-checked shorthand for node.state[pos].
    """
    return self.state[pos] if 0 <= pos < 49 else -1

```

2. `is_validMove(self, oldPos, direction):`

```

def is_validMove(self, oldPos, direction):
    """
    Return whether it is valid to move a peg from a given position
    in a given direction. The position must have a peg, the
    destination must be empty, and the intermediate position must
    have another peg.
    """
    midPos = self.getNextPosition(oldPos, direction)
    newPos = self.getNextPosition(midPos, direction)
    # 1 is a peg, and 0 is a hole
    # Since oldPos is valid, don't use bounds-checked indexing for it
    return (self.state[oldPos] == self[midPos] == 1
            and self[newPos] == 0)

```

3. `getNextPosition(self, oldPos, direction):`

```

@staticmethod
def getNextPosition(oldPos, direction):
    """
    Return the new position after moving in the given direction
    (-7 is north/up, 1 is east/right, 7 is south/down, and -1 is
    west/left). Invalid movements outside the bounds of the board
    will return positions such that is_corner(position) returns True.
    """
    newPos = oldPos + direction
    # If changing rows (north or south), newPos is already correct
    return (newPos if direction == 7 or direction == -7 or
            # If changing columns (east or west), the row should not change
            newPos // 7 == oldPos // 7 else -1)

```

4. `getNextState(self, oldPos, direction):`

```

def getNextState(self, oldPos, direction, pegSol):
    """
    Return a child node of the current one, created by a given valid
    move. The given game has its count of expanded nodes incremented.
    """
    pegSol.nodesExpanded += 1
    if not self.is_validMove(oldPos, direction):
        print "Error, You are not checking for valid move"
        exit(0)

```

```

midPos = self.getNextPosition(oldPos, direction)
newPos = self.getNextPosition(midPos, direction)
# x[:] makes a copy of x (necessary to avoid mutating
# self.state when updating childState)
childState = self.state[:]
childState[oldPos] = 0 # The peg moves from here, leaving a hole
childState[midPos] = 0 # The jumped-over peg is removed
childState[newPos] = 1 # The peg moves to this hole
# Convert positions back into pairs for printing
childTrace = self.trace + [(oldPos // 7, oldPos % 7),
                           (newPos // 7, newPos % 7)]
return gameNode(childState, childTrace,
                self.pegCount - 1, self.heuristic)

```

2.3 Search functions

We were required to implement our search functions and heuristics in *search.py*.

We implemented ID-DFS straightforwardly, based on figure 3.17 (section 3.4, page 88) and figure 3.18 (section 3.4, page 89) in the textbook. However, we also maintain a set of explored nodes (modulo symmetry) to avoid revisiting, which helps prune the search tree. Nodes for which a depth-limited search fails to find a solution (because one does not exist, not because it was cut off before finding one) are kept in a second explored set which is reused for the even deeper searches.

We implemented A* based on figure 3.14 (section 3.4, page 84) in the textbook. However, we take the set of *explored* nodes modulo symmetry to further prune the search tree. (One game state can be solved if and only if its seven symmetrical states can be solved, so when one has been expanded the other seven can be pruned.)

2.3.1 First A* heuristic

```

def heuristicOne(node):
    """
    This heuristic counts the number of "dangling" pegs (those without
    any neighbors) and adds it to twice the number of pegs remaining
    on the board.

    We penalize dangling pegs because they cannot be removed without
    first moving another peg to a neighboring hole, which may not be
    possible.
    """
    num_dangling = 0
    for i in xrange(49):
        if node.state[i] != 1: continue
        for direction in [7, 1, -7, -1]:
            if node[i + direction] == 1:
                break
        else:
            num_dangling += 1
    return node.pegCount * 2 + num_dangling

```

2.3.2 Second A* heuristic

```

def heuristicTwo(node):
    """

```

This heuristic sums the estimated difficulty of removing each peg and adds it to twice the number of pegs remaining on the board.

We chose the difficulty values after some trial and error. The outer four areas of the board are generally less maneuverable than the center, and further from the eventual goal of the very central hole. This applies especially to the eight outer corners, which have fewer neighbors to jump over. The four inner corners are only slightly easier, since they have more neighbors but are still not in the ideal rows or columns. (Note that pegs in the outer and inner corners can only move within those positions, never to a non-corner hole.) The center peg is the end goal of the game, but prior to that having a peg sit there is not quite useful (since if you are left with two pegs and one is in the center, the game is unsolvable).

```
"""
difficulties = [
    0, 0, 4, 1, 4, 0, 0,
    0, 0, 1, 1, 1, 0, 0,
    4, 1, 2, 0, 2, 1, 4,
    1, 1, 0, 1, 0, 1, 1,
    4, 1, 2, 0, 2, 1, 4,
    0, 0, 1, 1, 1, 0, 0,
    0, 0, 4, 1, 4, 0, 0,
]
return node.pegCount * 2 + sum(difficulties[k]
    for k in xrange(49) if node.state[k] == 1)
```

2.3.3 Other A* heuristics

We tested various other heuristics before settling on our chosen two. One of them was a baseline heuristic that simply returned the number of remaining pegs. The other three were all variations on *heuristicTwo* with a different *difficulties* array.

1. Manhattan distance from center neighbors:

```
distances = [
    5, 4, 3, 2, 3, 4, 5,
    4, 3, 2, 1, 2, 3, 4,
    3, 2, 1, 0, 1, 2, 3,
    2, 1, 0, 1, 0, 1, 2,
    3, 2, 1, 0, 1, 2, 3,
    4, 3, 2, 1, 2, 3, 4,
    5, 4, 3, 2, 3, 4, 5
]
```

2. Alternative difficulties #1:

```
difficulties = [
    0, 0, 4, 0, 4, 0, 0,
    0, 0, 0, 0, 0, 0, 0,
    4, 0, 3, 0, 3, 0, 4,
    0, 0, 0, 1, 0, 0, 0,
    4, 0, 3, 0, 3, 0, 4,
    0, 0, 0, 0, 0, 0, 0,
]
```

```

    0, 0, 4, 0, 4, 0, 0
]

```

3. Alternative difficulties #2:

```

difficulties = [
    0, 0, 1, 1, 1, 0, 0,
    0, 0, 1, 1, 1, 0, 0,
    1, 1, 0, 0, 0, 1, 1,
    1, 1, 0, 0, 0, 1, 1,
    1, 1, 0, 0, 0, 1, 1,
    0, 0, 1, 1, 1, 0, 0,
    0, 0, 1, 1, 1, 0, 0
]

```

3 Results

3.1 Test cases

Our first test case was the given *game.txt*, a board with 6 pegs in a cross shape.

We wrote ten more test cases manually based on actual game boards that people play, such as a 9-peg plus shape, a 16-peg pyramid, a 24-peg diamond, or the 32-peg “central game” with only one hole in the center.

We also created 100 random test cases with between 2 and 21 pegs, by implementing a generator for random valid test cases. It works by starting with the goal state (only one peg in the center) and randomly picking valid backward moves that could have reached that state. In pseudocode:

Algorithm 1 Random test case generator

```

1: numSteps ← pick a random number between 1 and 33
2: boardState ← goal board state, with only one peg in the center
3: for j ← 1 to numSteps do
4:   validStartPos, validDirection ← pick a random valid move
5:   boardState[validStartPos] ← 0
6:   boardState[validStartPos + validDirection] ← 1
7:   boardState[validStartPos + validDirection + validDirection] ← 1
8: end for
9: return board_state

```

3.2 Benchmarks

We compared the performance of our different search algorithms on a few test cases:

1. The given 6-peg *game.txt*
2. The 32-peg central game
3. The mean expanded node count of all 110 test boards
4. The mean count without the central game outlier
5. The median count of all 110 test boards

Test case	ID-DFS	A*1	A*2	A*Base	A*ManD	A*Alt1	A*Alt2
game.txt	50	16	13	13	19	13	19
Central game	MemoryError	4,713,300	103,402	26,807	1,801,655	170	530,402
Mean	1,979,736	47,084	3,388	5,146	19,201	4,394	8,917
Mean w/o central	1,979,736	4,275	2,470	4,948	2,848	4,433	4,133
Median	6,417	74	98	92	95	128	114

Table 1: Expanded node counts for all algorithms in different test cases

The performance results for our different search algorithms are shown in table 1:

We also tested each algorithm on a board with no solution (two noncontiguous pegs). They all succeeded in printing “Impossible to solve” instead of a solution move trace, without crashing or infinitely looping.

The heuristics A*2, A*ManD, A*Alt1, and A*Alt2 were conceptually very similar, so we only wanted to submit one of them. Our tests found that A*2 was the most efficient: it had the lowest mean count (of all the algorithms), and while A*Manhattan had a slightly lower median count, it could not solve the central game quickly enough.

We never intended to submit A*Base as a heuristic, since it was so simple, so we were surprised to see it perform as well as it did. It had a competitive mean expanded node count and the lowest median count. We hypothesize that our more complicated heuristics are better for particular game boards, but not for every possible random board that we were trying.

3.3 Performance notes

In our testing, we found an interesting phenomenon. When enumerating valid moves from a particular position in a particular game state, there can be at most four moves (north, south, east, and west). For several large test cases, if we test those four directions in a different order, performance can significantly vary. There are 4!, or 24 possible orders, so we tested all 24 of them against all six search algorithms. The 24 possibilities break down into two categories: “opposite directions first” (such as “north, south, east, west,” where “south” is opposite “north”) and “orthogonal directions first” (such as “north, east, south, west,” where “east” is orthogonal to “north”). By plotting the results in figure 1, it is clear that the “orthogonal directions first” category had slightly better average performance, but one outlier that did noticeably worse. We chose to test the directions orthogonal-first (specifically “south, east, north, west”).

4 Conclusion

We fulfilled the requirements to implement the three search algorithms, changing only the parts of the code in specified locations, with appropriate comments. Our algorithms correctly handle all test cases, including impossible boards, and they perform well even on large boards with the help of pruning symmetric states.




Figure 1: Box plot for two categories of direction order